

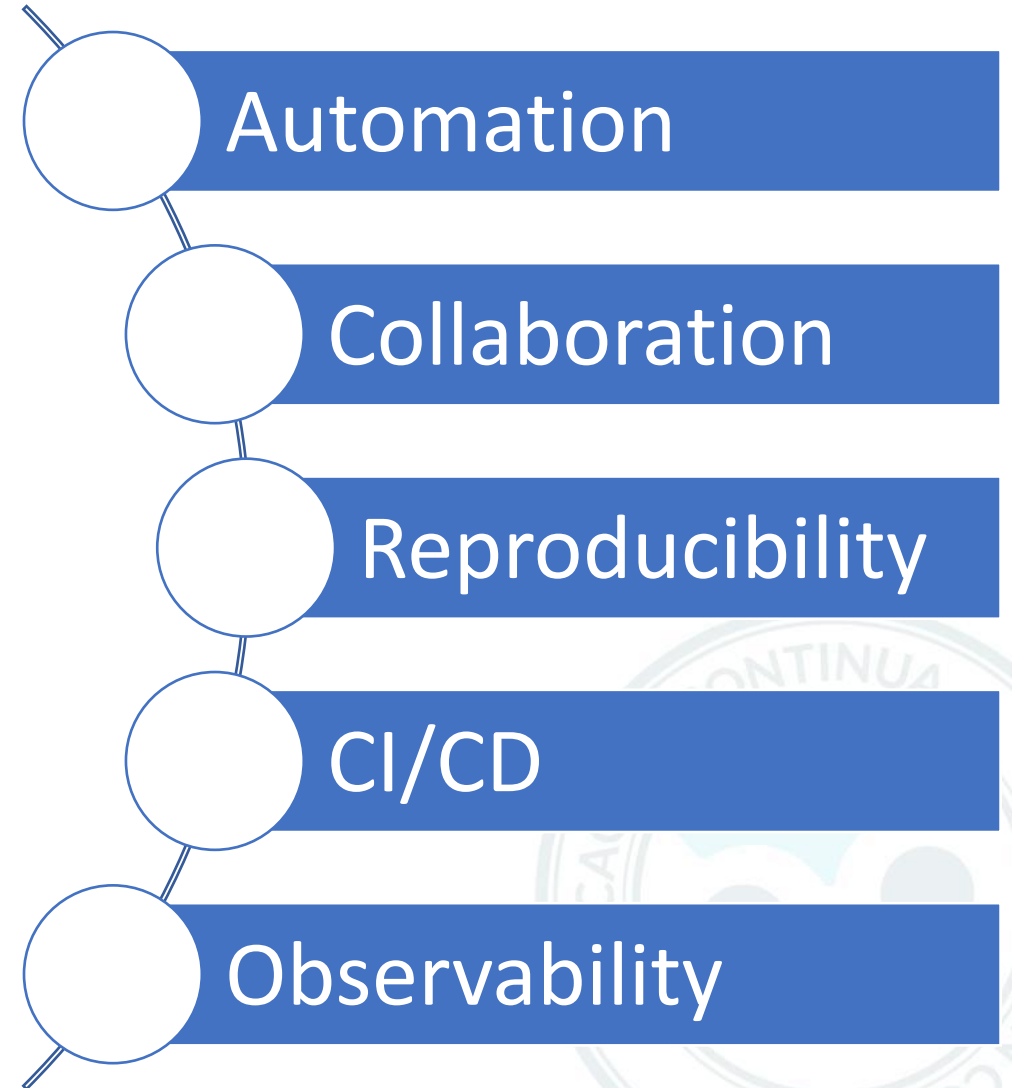

Large Language Models Operations (LLMOps)




- LLMOPs introduction
- Deployment and scalability of LLMs
- Monitoring and maintenance of models in production
- Performance evaluation and continuous improvement
- Ethical considerations and privacy



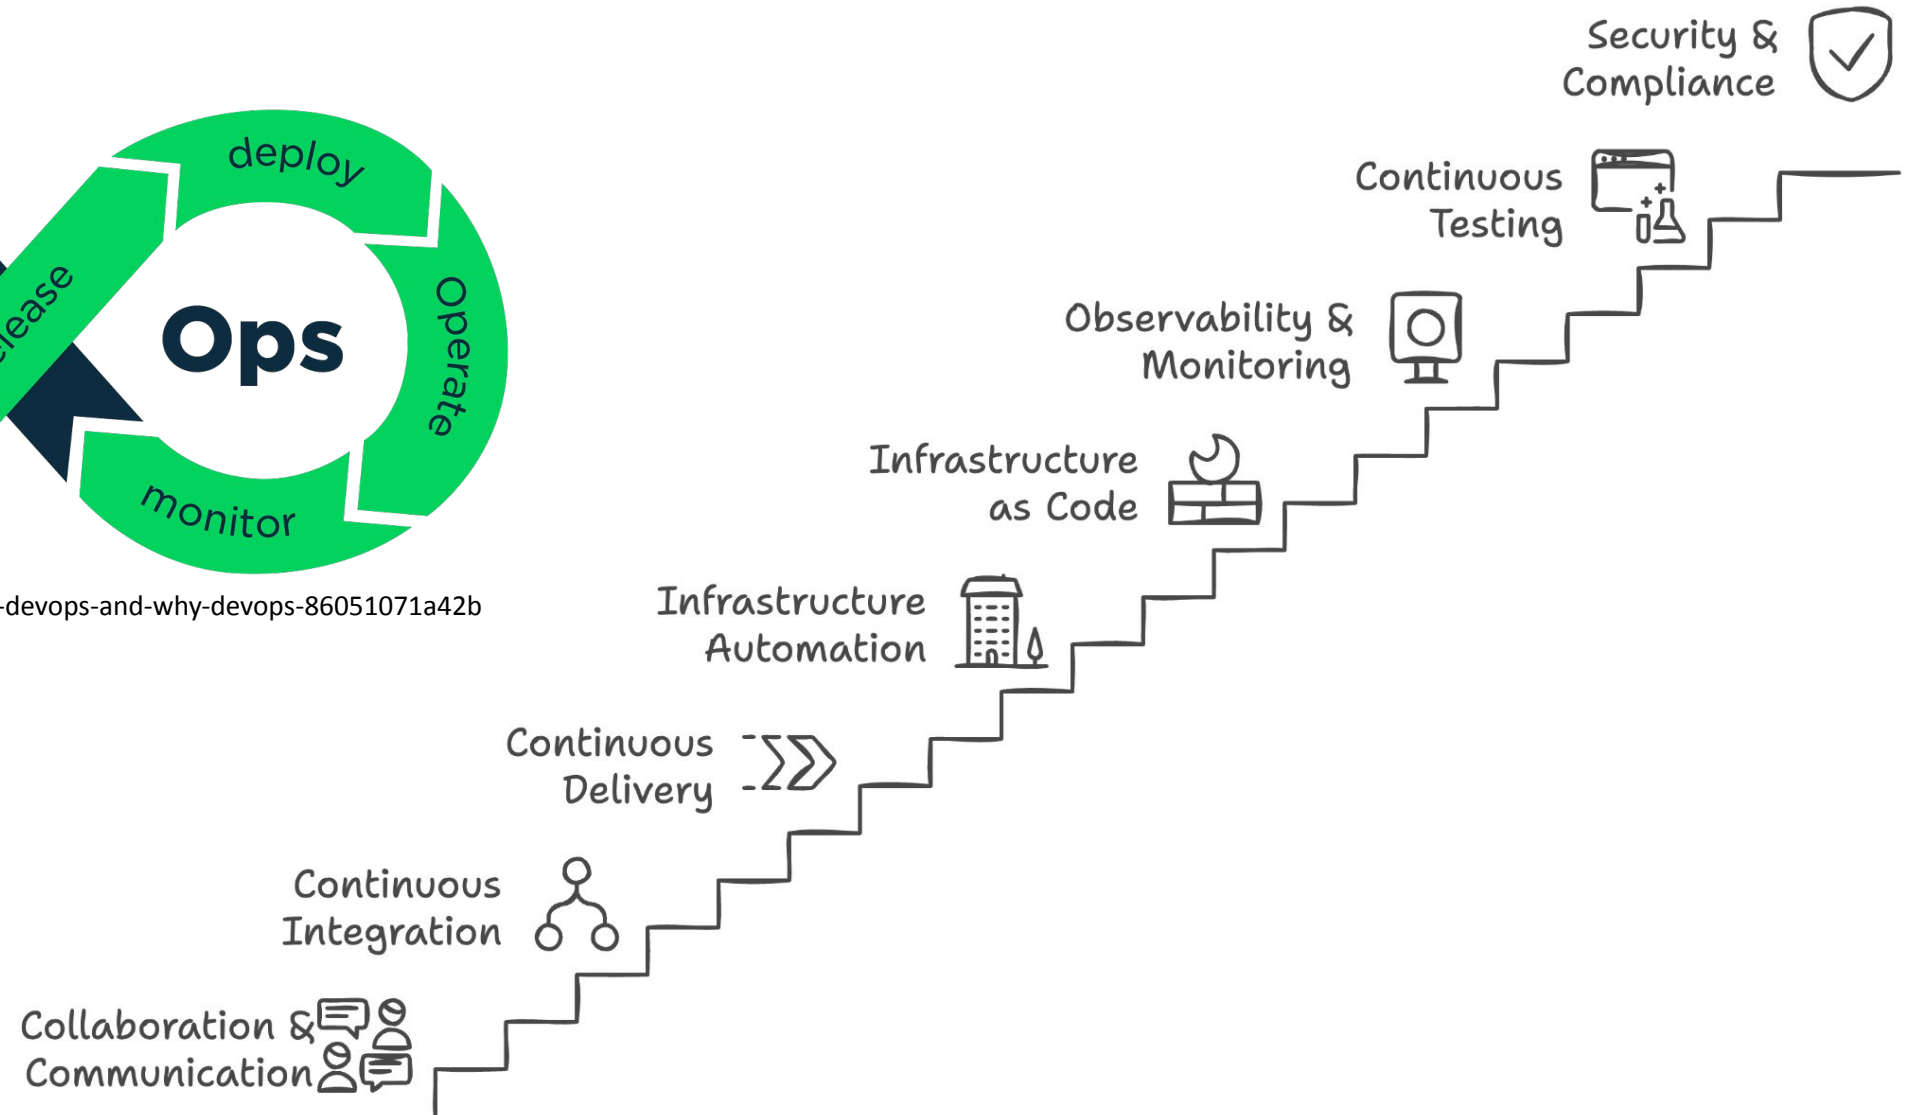
Intro to LLMOPS

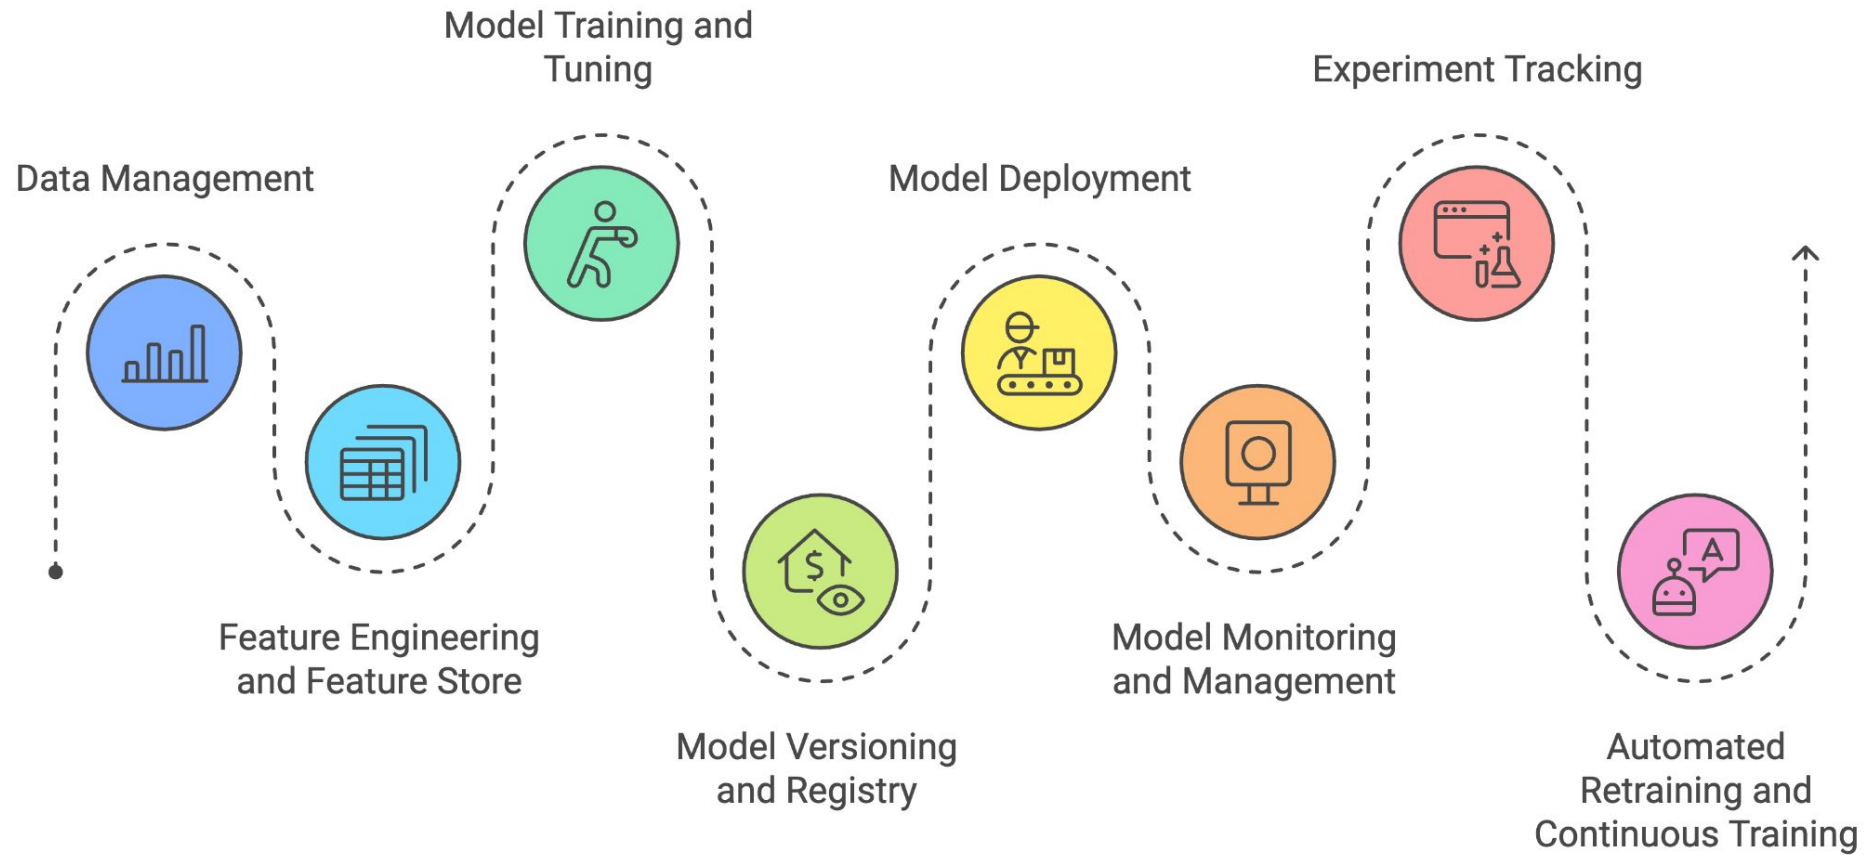


DevOps



<https://medium.com/@rituserke86/what-is-devops-and-why-devops-86051071a42b>





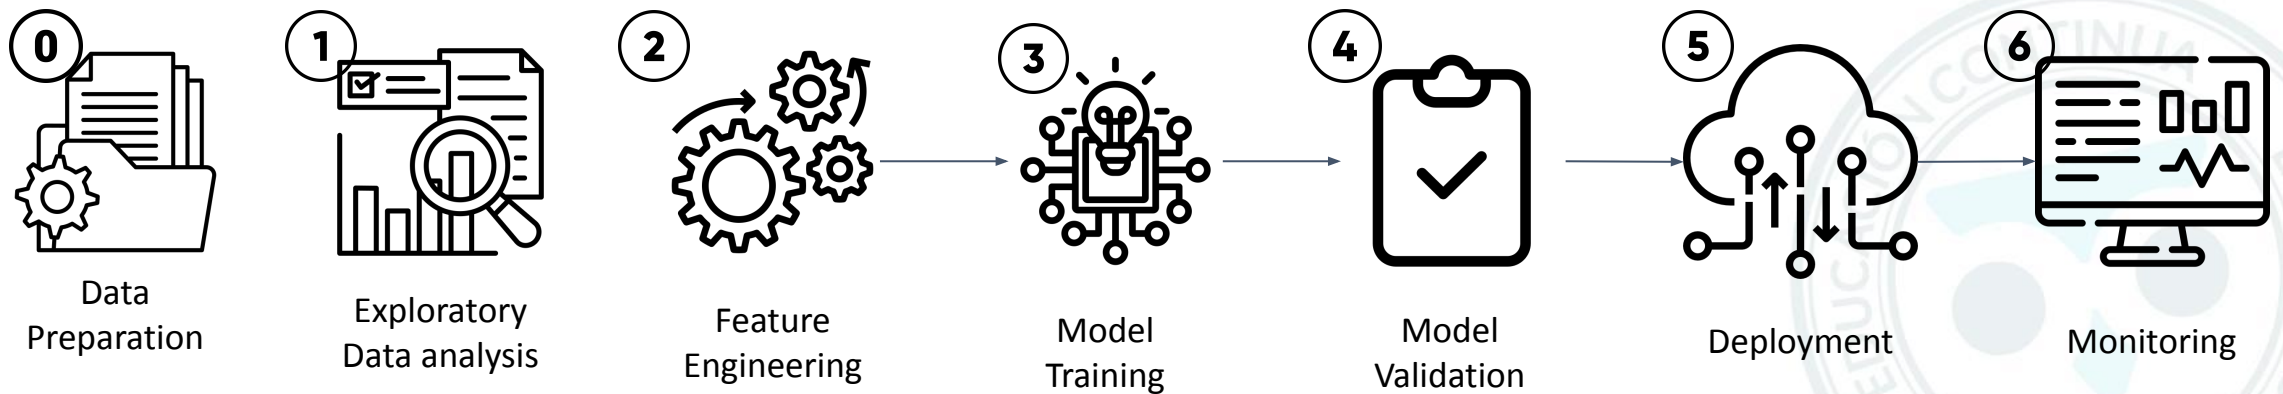
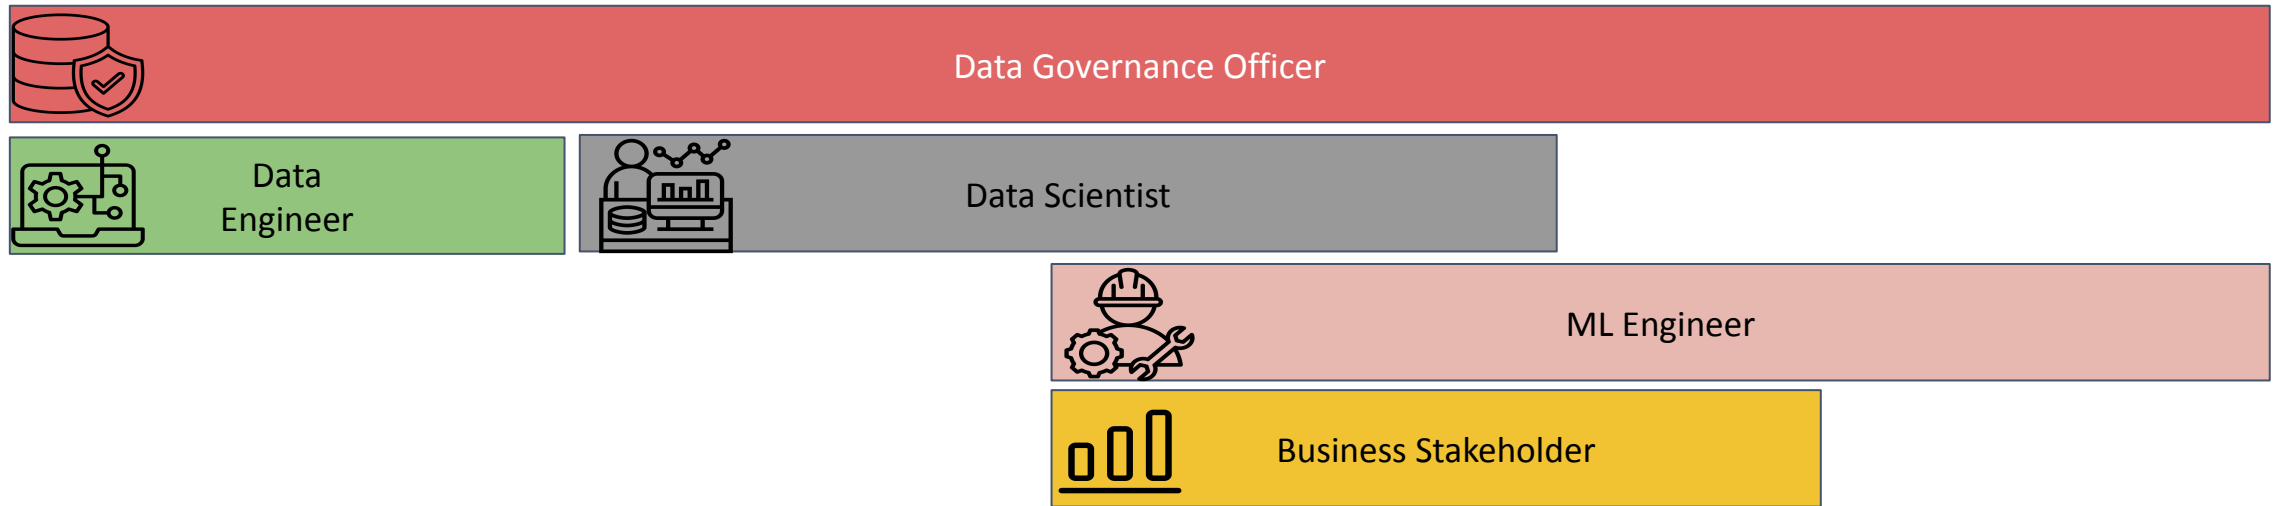
LLMOps



Training and deployment of Large Language Models



MLOps



What specific scaling challenges exist for LLMs?



Initial training: trillions of tokens, hundreds to thousands of GBs & very long run times.



Fine-tuning: updating model weights based on your own data, still requires relatively large data and long training times. Plus lots of evaluation!



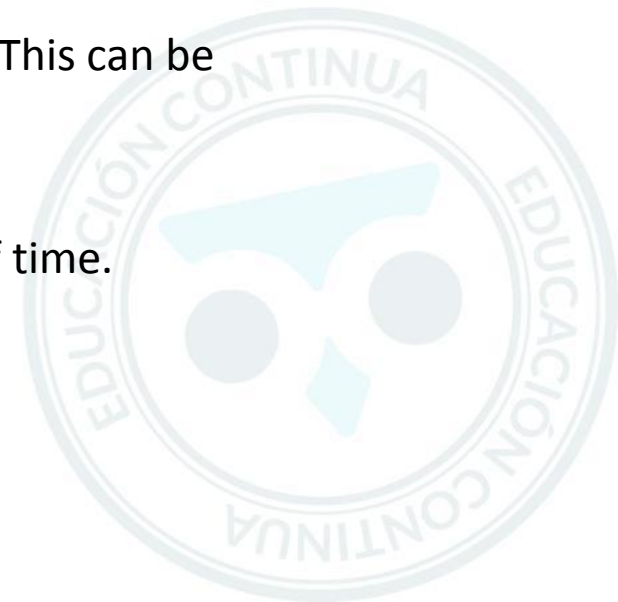
Storage: the models contain billions of parameters, often hundreds of GB. This can be prohibitive for some devices.



Latency: Running inference on these models can be very costly in terms of time.



Cost: All of this costs \$\$\$!



What specific scaling challenges exist for LLMs?



Initial training: don't do it!



Fine-tuning: Optimize and use a scalable framework, such as Ray.



Storage: Quantization, memorization, caching ...



Latency: Quantization, memorization, caching, hardware and memory bandwidth optimization...



Cost: Above plus use 'open source' models



Types of use-cases for Enterprises

How willing are enterprises to use LLMs for different use cases?



(% of enterprises experimenting with given use case who have deployed to production)



Source: a16z survey of 70 enterprise AI decision makers

LLM application archetypes

Prompt
Engineering

Retrieval
Augmented
Generation



Agents

Multi-Agents


Fine-tunning



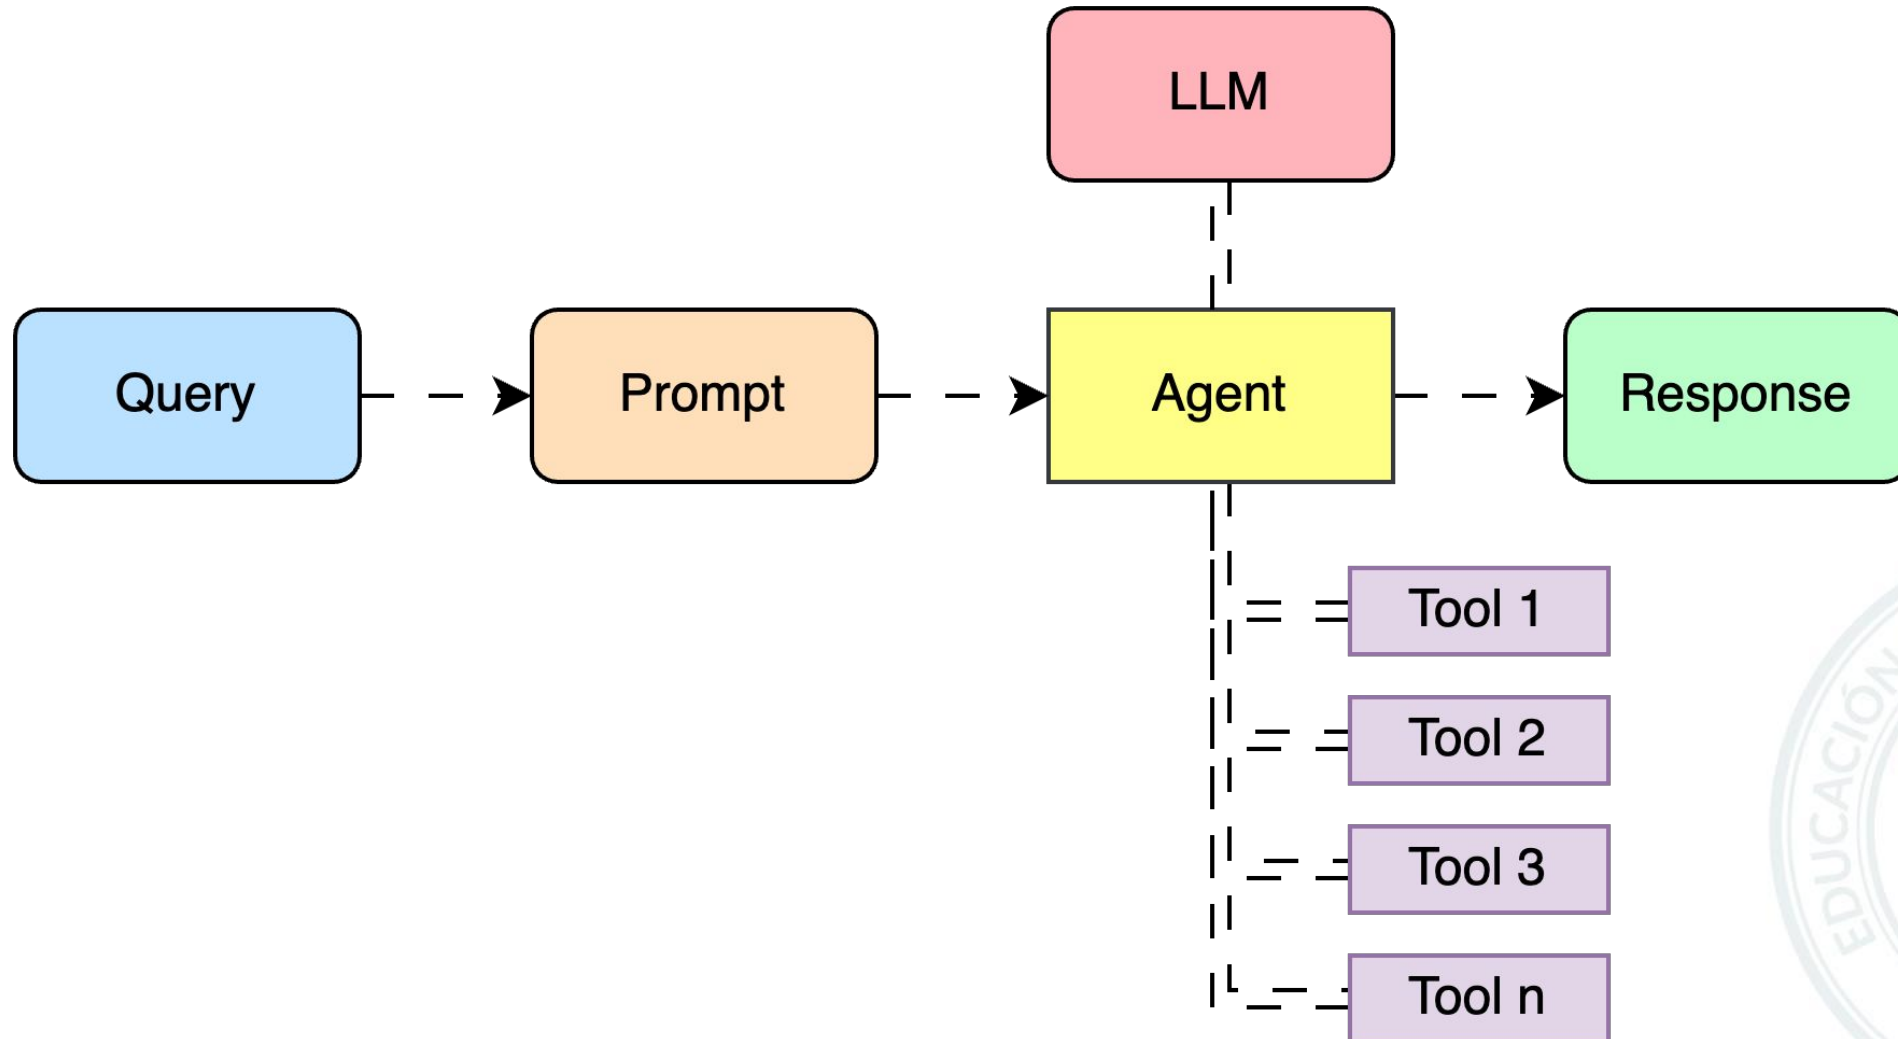
Prompt Engineering Applications



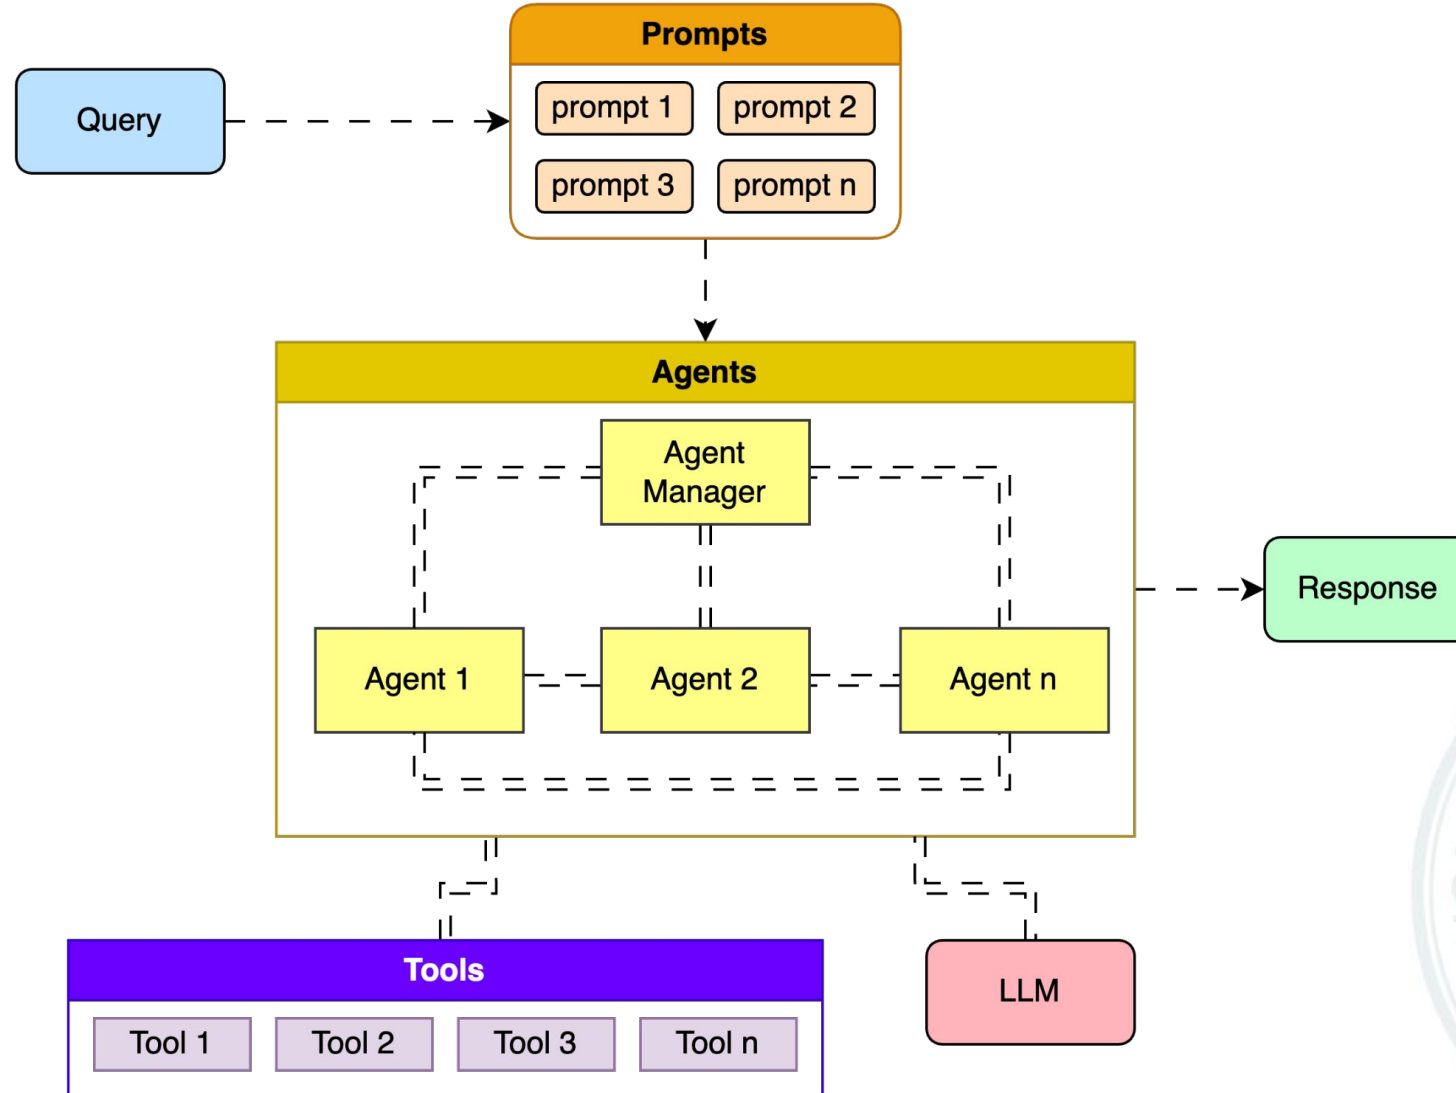
Retrieval Augmented Generation Applications



Agentic Applications

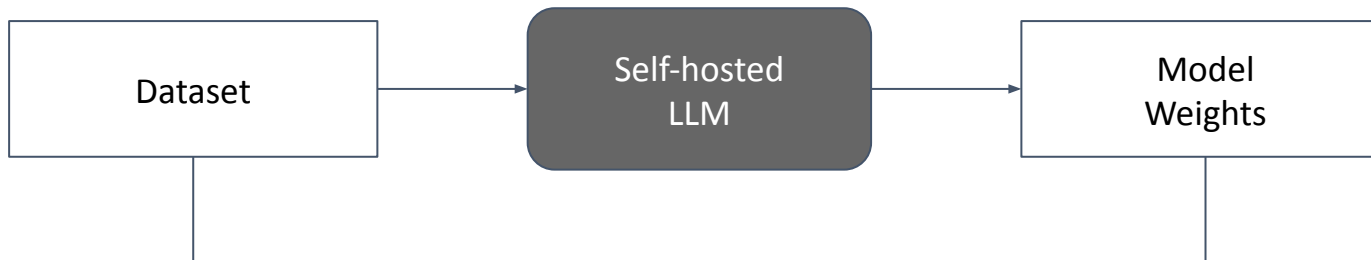


Multi-Agent Applications



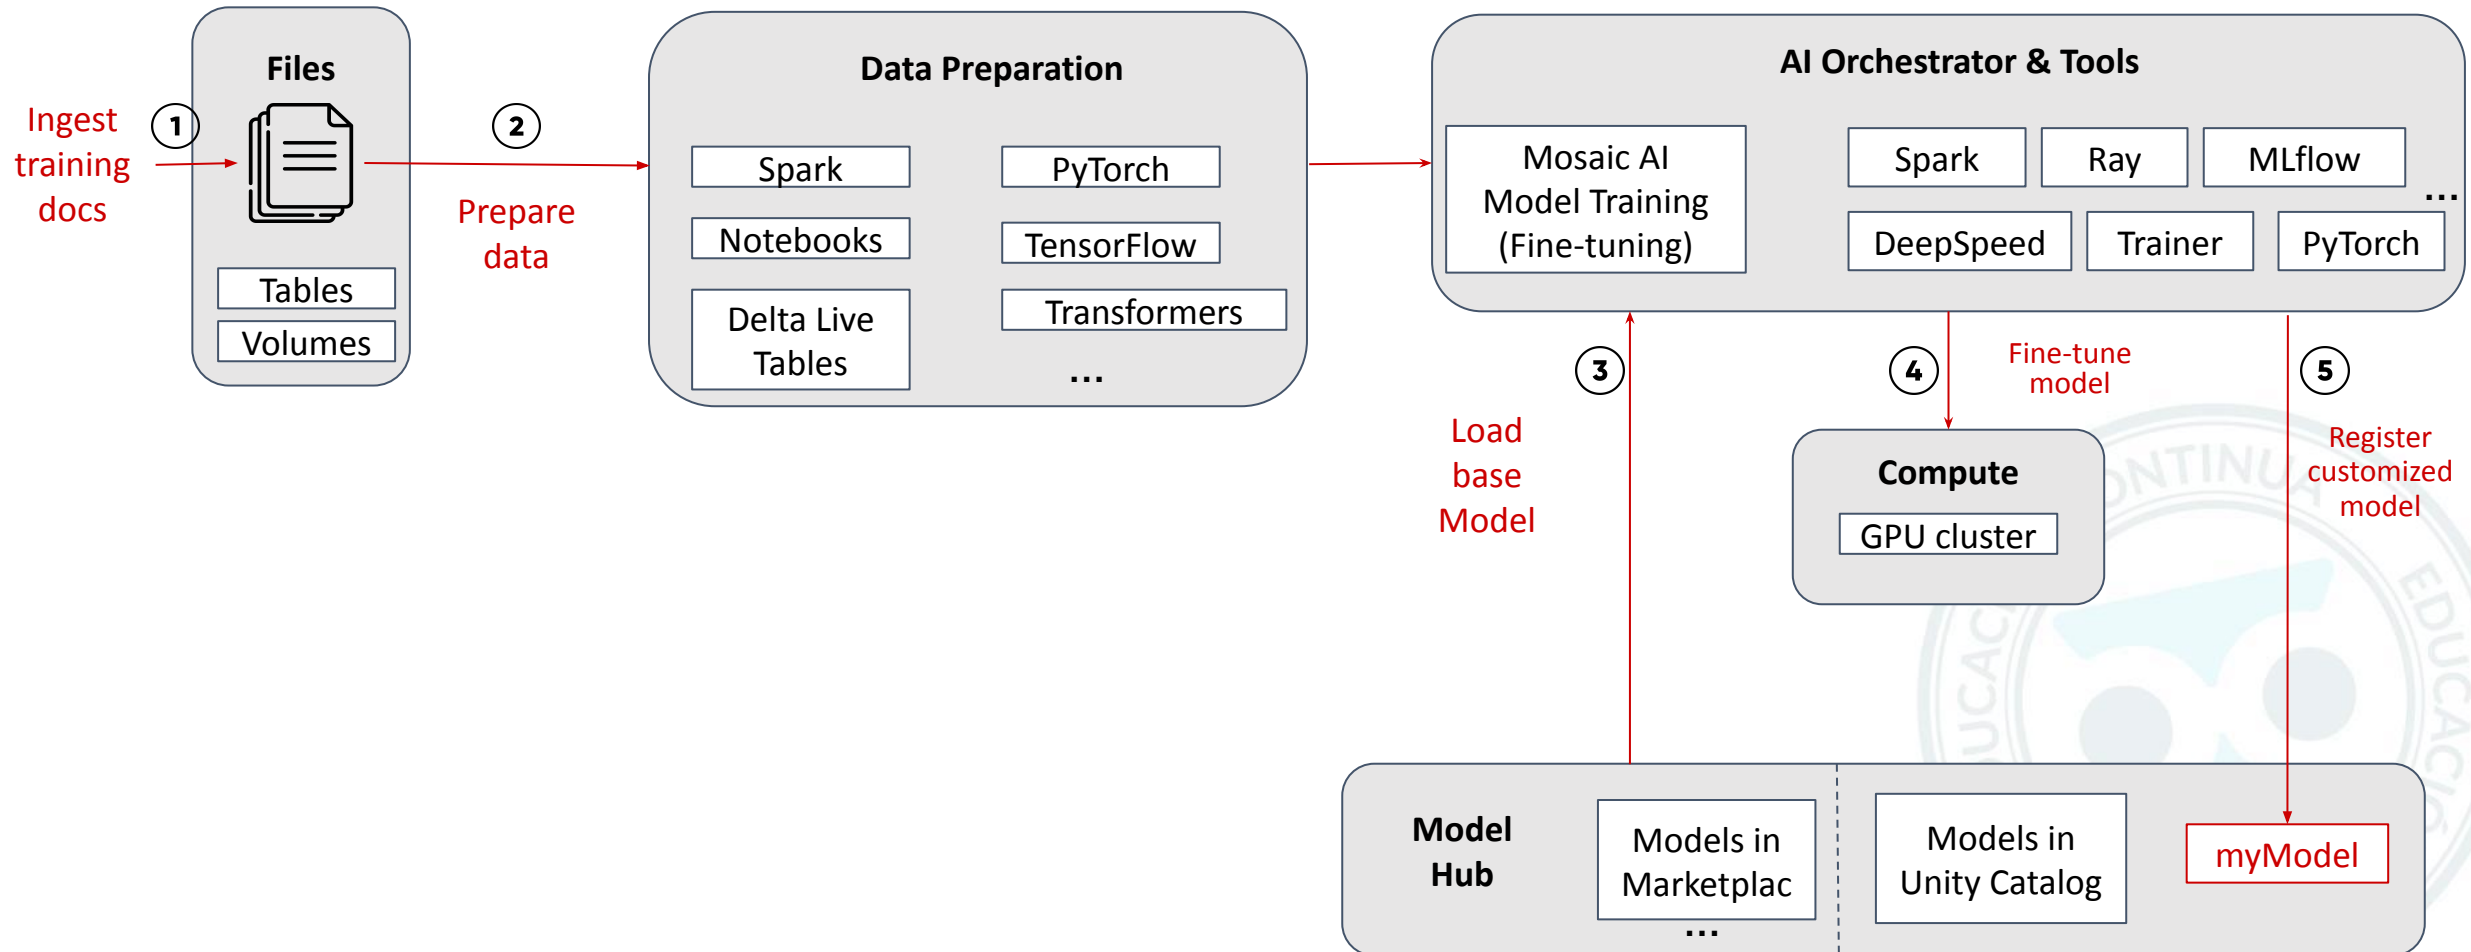
Operationalizing Fine-tuning Pipelines

- Automation, version control, reproducibility
- Distributed training infrastructure
 - DeepSpeed, PEFT, GPUs
- Similar to classical model training serving
- Optimization techniques
 - vLLM, MLC, CudaGraph, MQA, Quantization, TensorRT
- Deeper understanding of GPUs, TTFT, TPOT



Types of LLM Applications

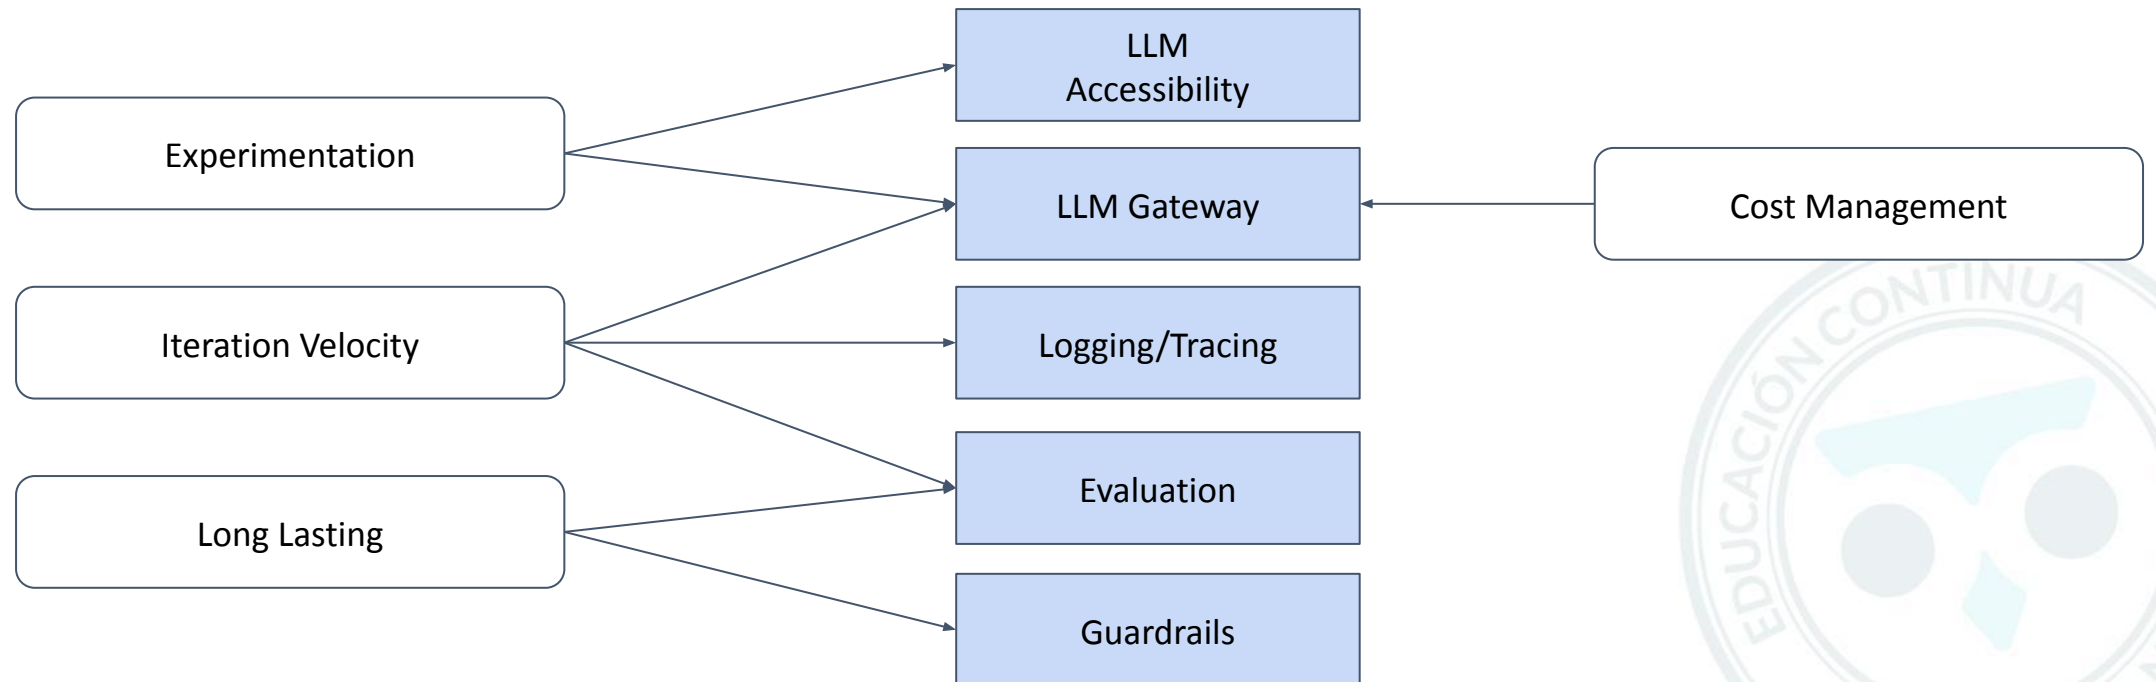
Architecture: fine-tuning



Types of LLM Applications

Practitioner's Perspective

LLMOps Starting Points



SUMMARIZATION TASKS FOR LLM

Improve asset management through LLM optimization and advanced summarization techniques

Key takeaways:

- Automation of summarization processes, facilitated by LLMs
- Data summarization extracts valuable insights, key trends, and strategic directives from large volumes of data to improve equipment management
- Evaluating and comparing LLM prompts optimizes performance and leads to more user-centric results
- Utilizing a weighted Quality Score model on the real project environment to improve operational efficiency and simplify decision-making

Authors:

Luis Angarita Gutiérrez
Constanza Garcia Diaz

softserve

Improve Asset Management with LLM Evaluation Pipelines

Creating efficient AI-powered summaries, LLMs save up to 93% of the time taken by traditional manual processes. Our white paper explains how this allows you to optimize asset management operations.

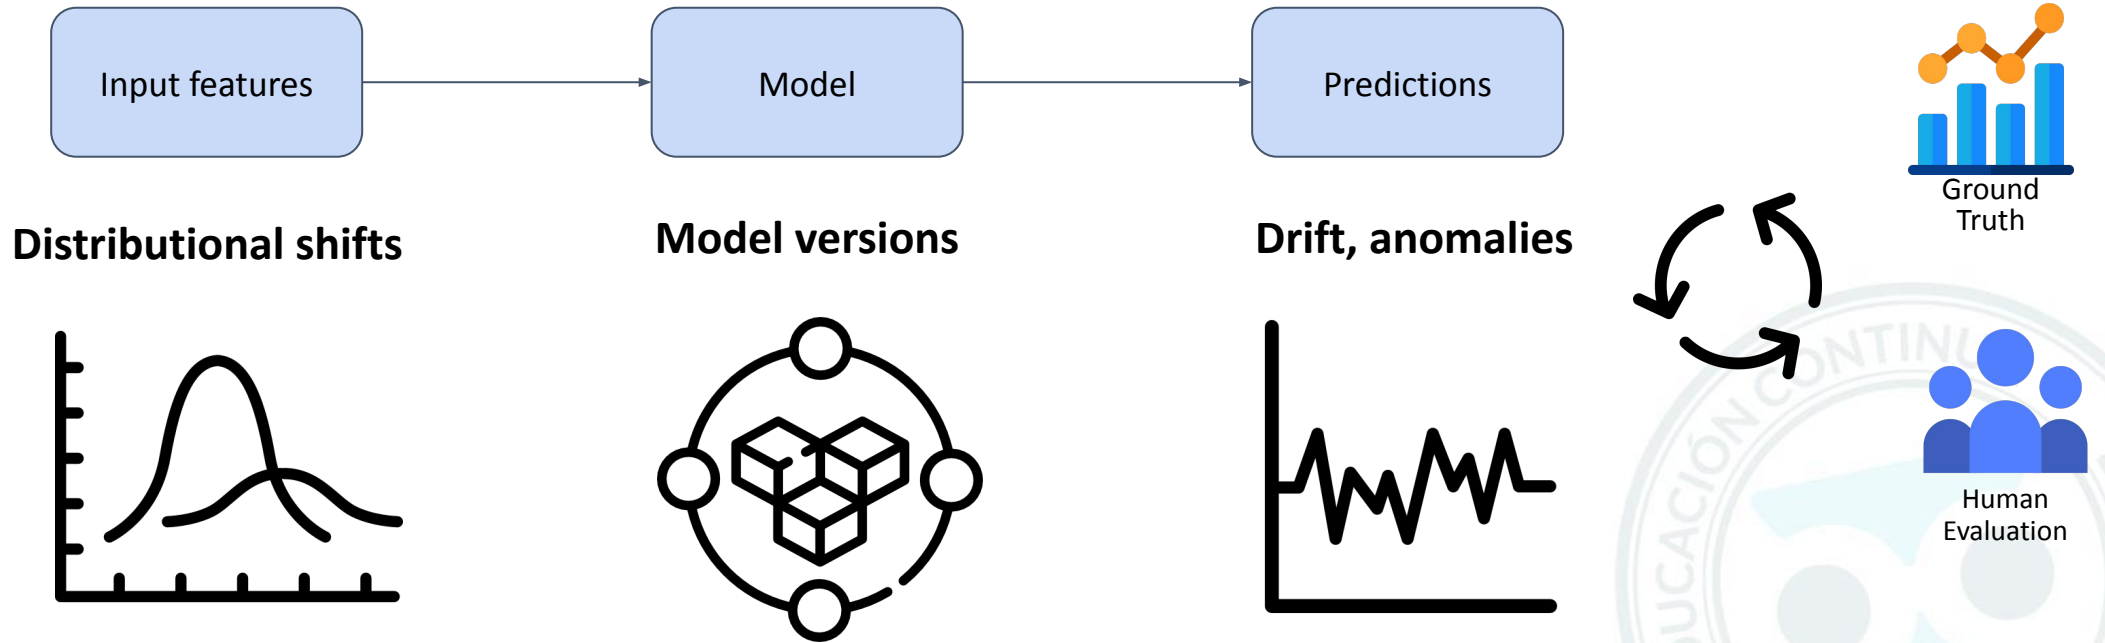
<https://info.softserveinc.com/summarization-tasks-for-llm-white-paper>



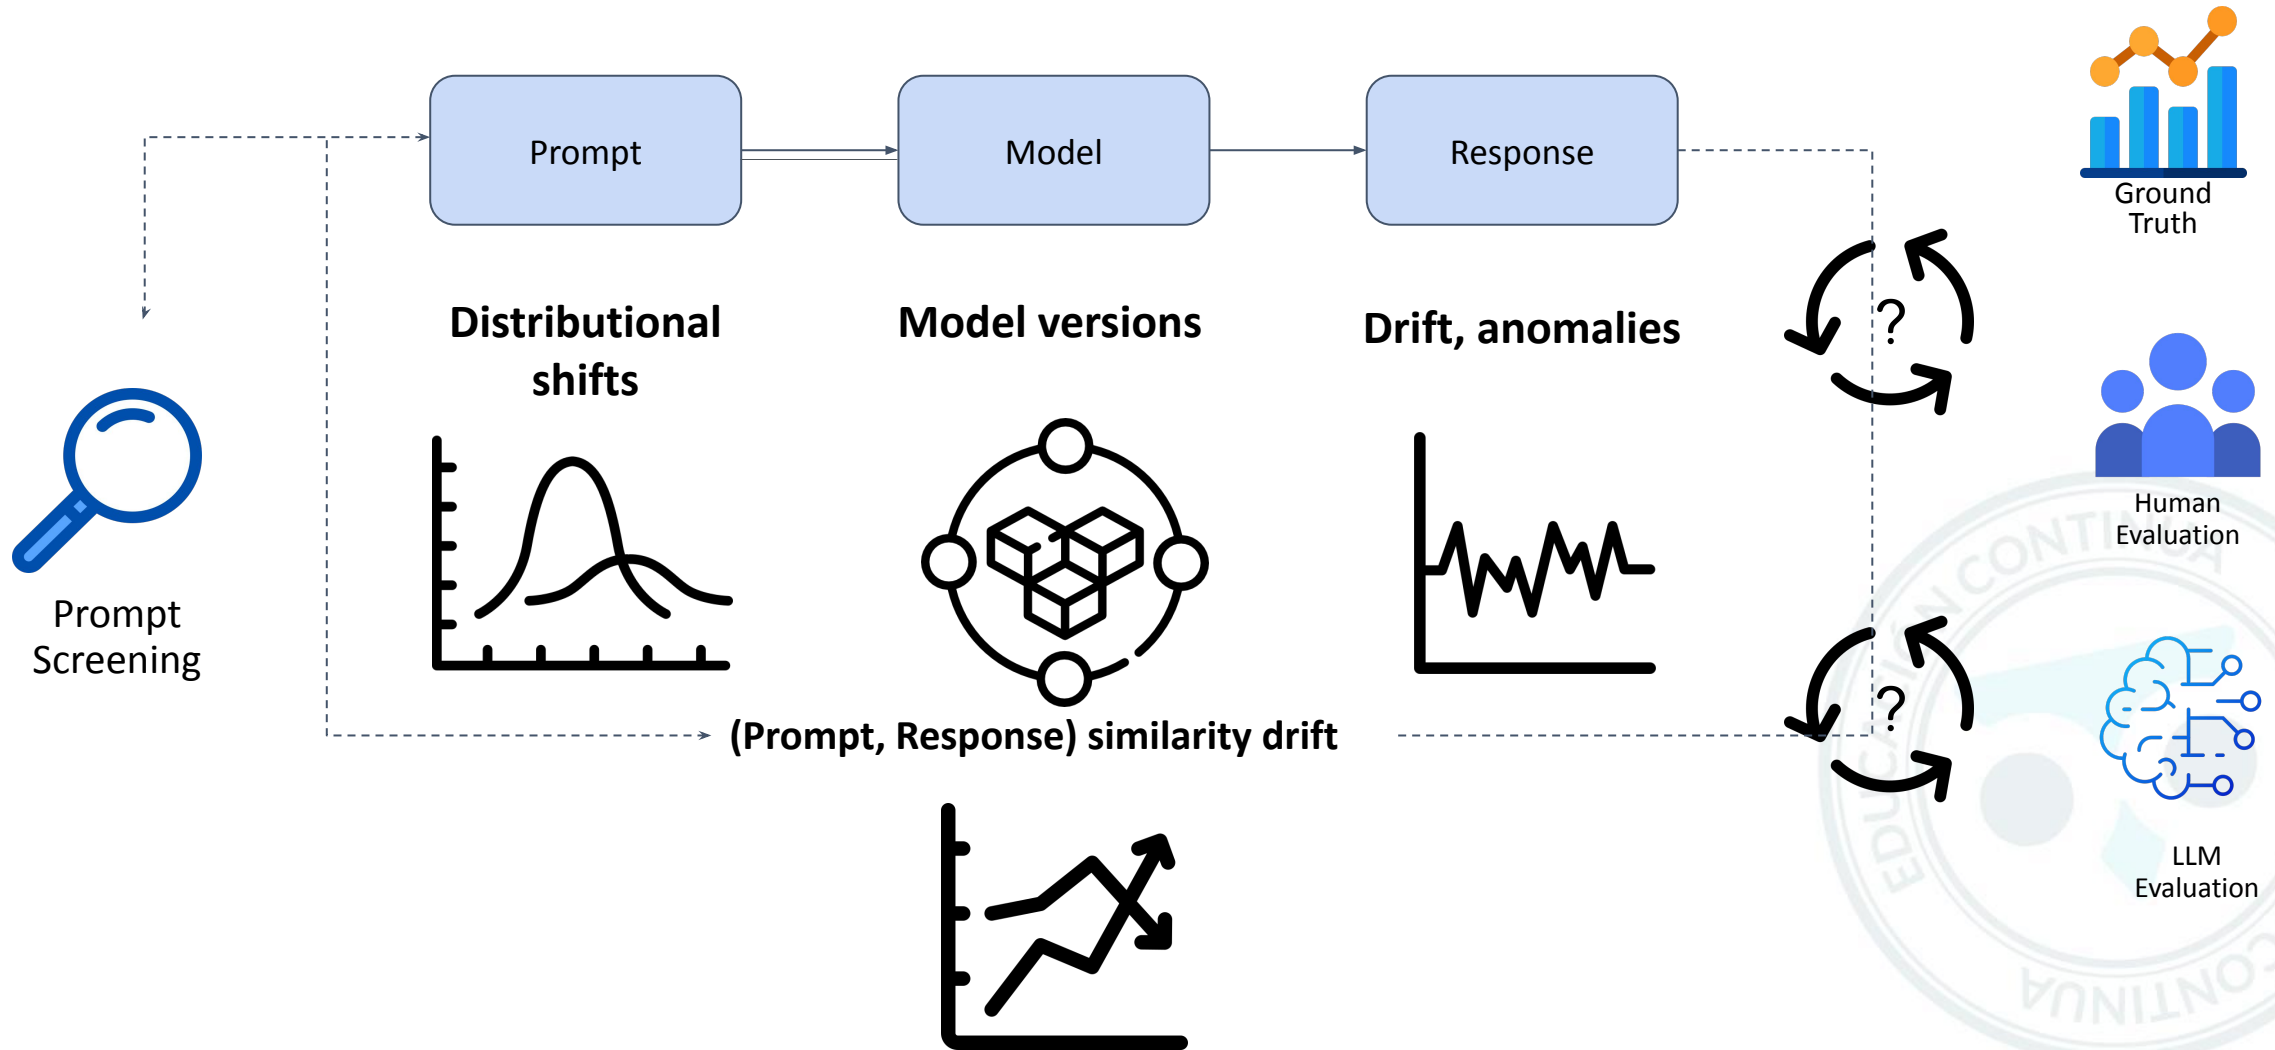
Monitoring and maintenance of LLM applications

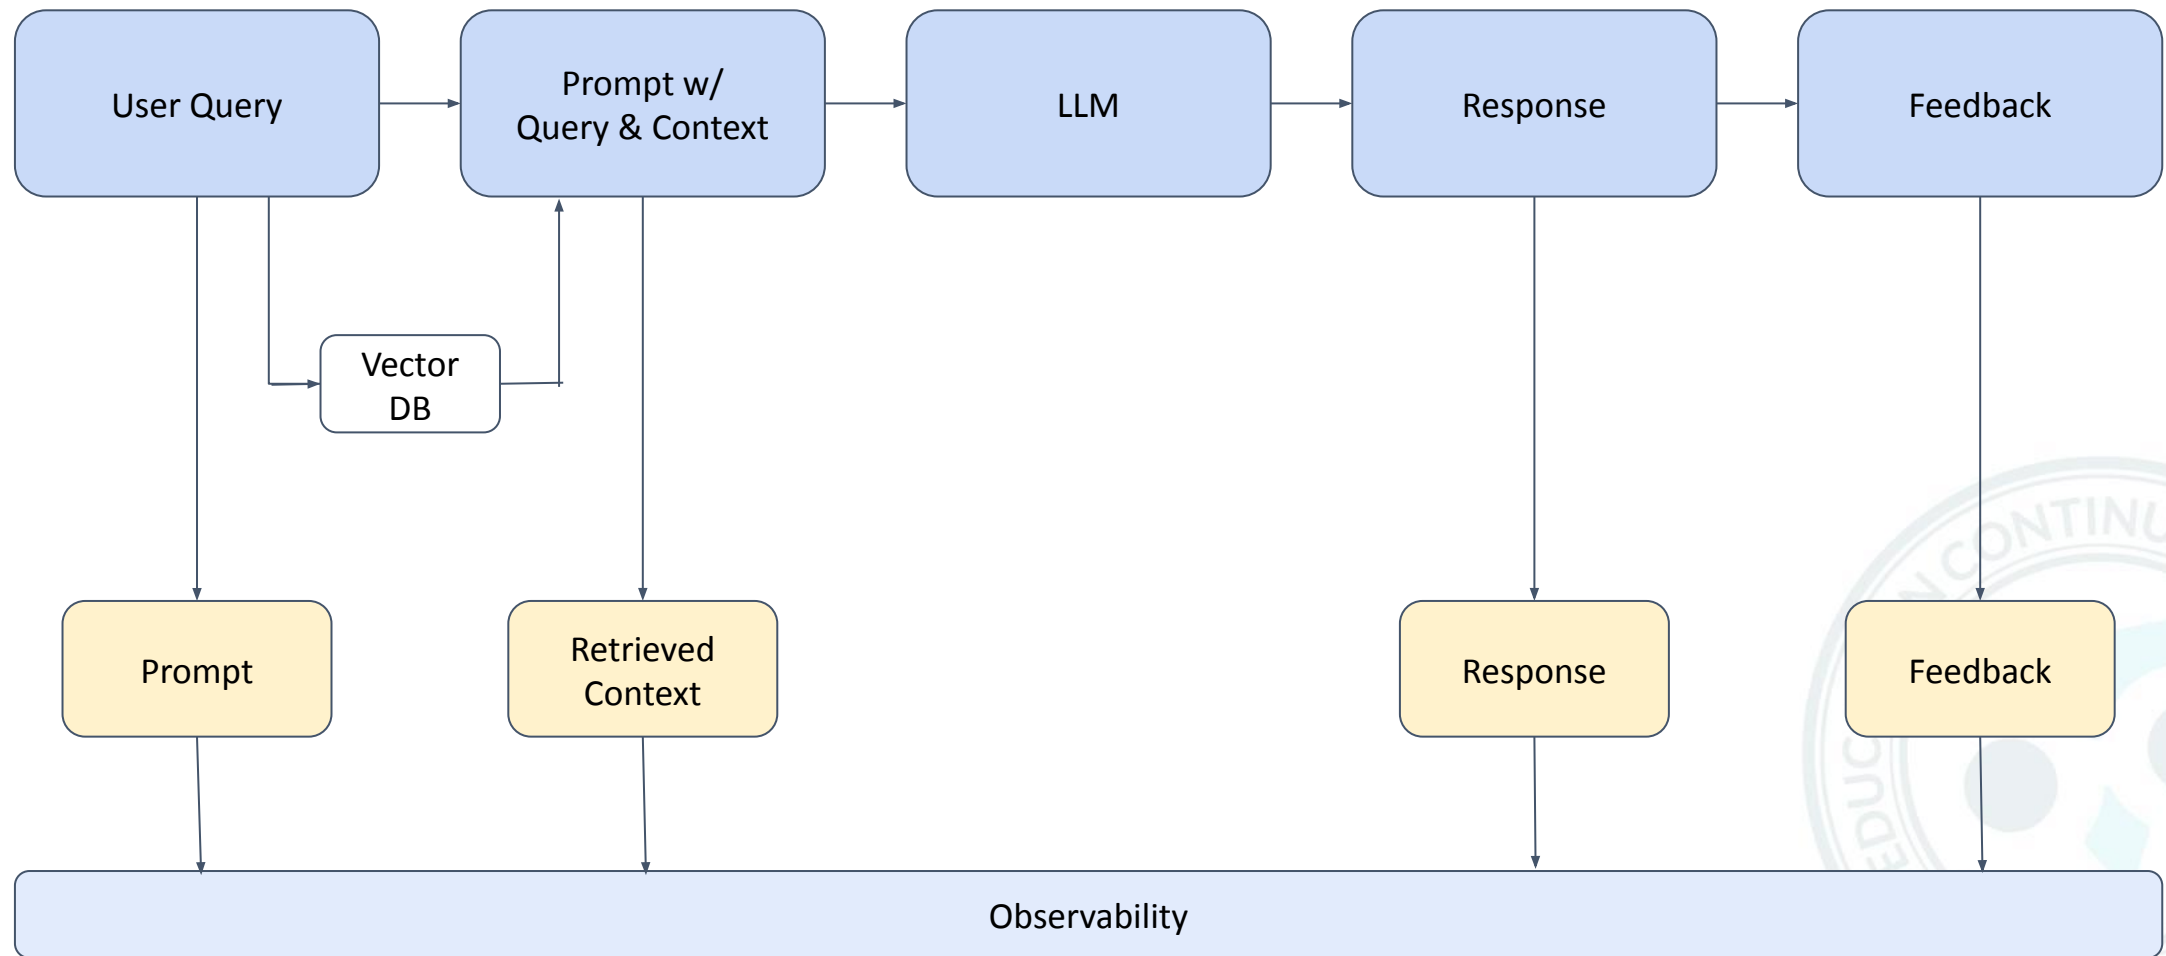


Monitoring for 'classic' ML



Monitoring for LLMs





Deploying and monitoring systems

Key advice

Use flexible tooling for packaging

Why?

- You will swap AI libraries over time: LangChain, LlamaIndex, Python, ...
- Uniform APIs lower the cost of switching libraries for a use case

How?

- MLflow supports built-in-flavors, PyFuncs, and custom flavors.
- All are managed behind uniform APIs

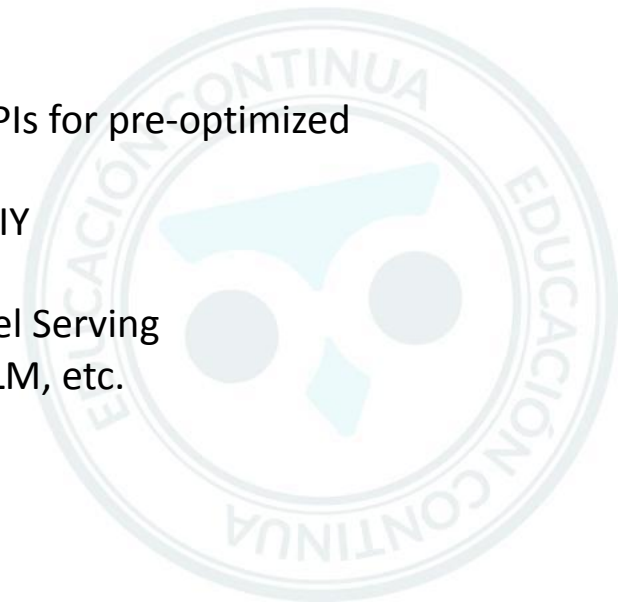
Use optimized inference

Why?

- User experience and TCO

How?

- Real-time: Model Serving
 - Foundation Model APIs for pre-optimized architectures
 - Custom models for DIY
- Batch and streaming
 - ai_query to call Model Serving
 - GPU clusters with vLLM, etc.



Performance evaluation and continuous improvement of LLM applications

LLM evaluation methods & resources

Ground-truth metrics




- BLEU
- ROGUE
- METEOR

Benchmarks




- TruthfulQA
- Arc
- HellaSwag

Uncertainty estimation



- SelfCheckGPT
- Perplexity

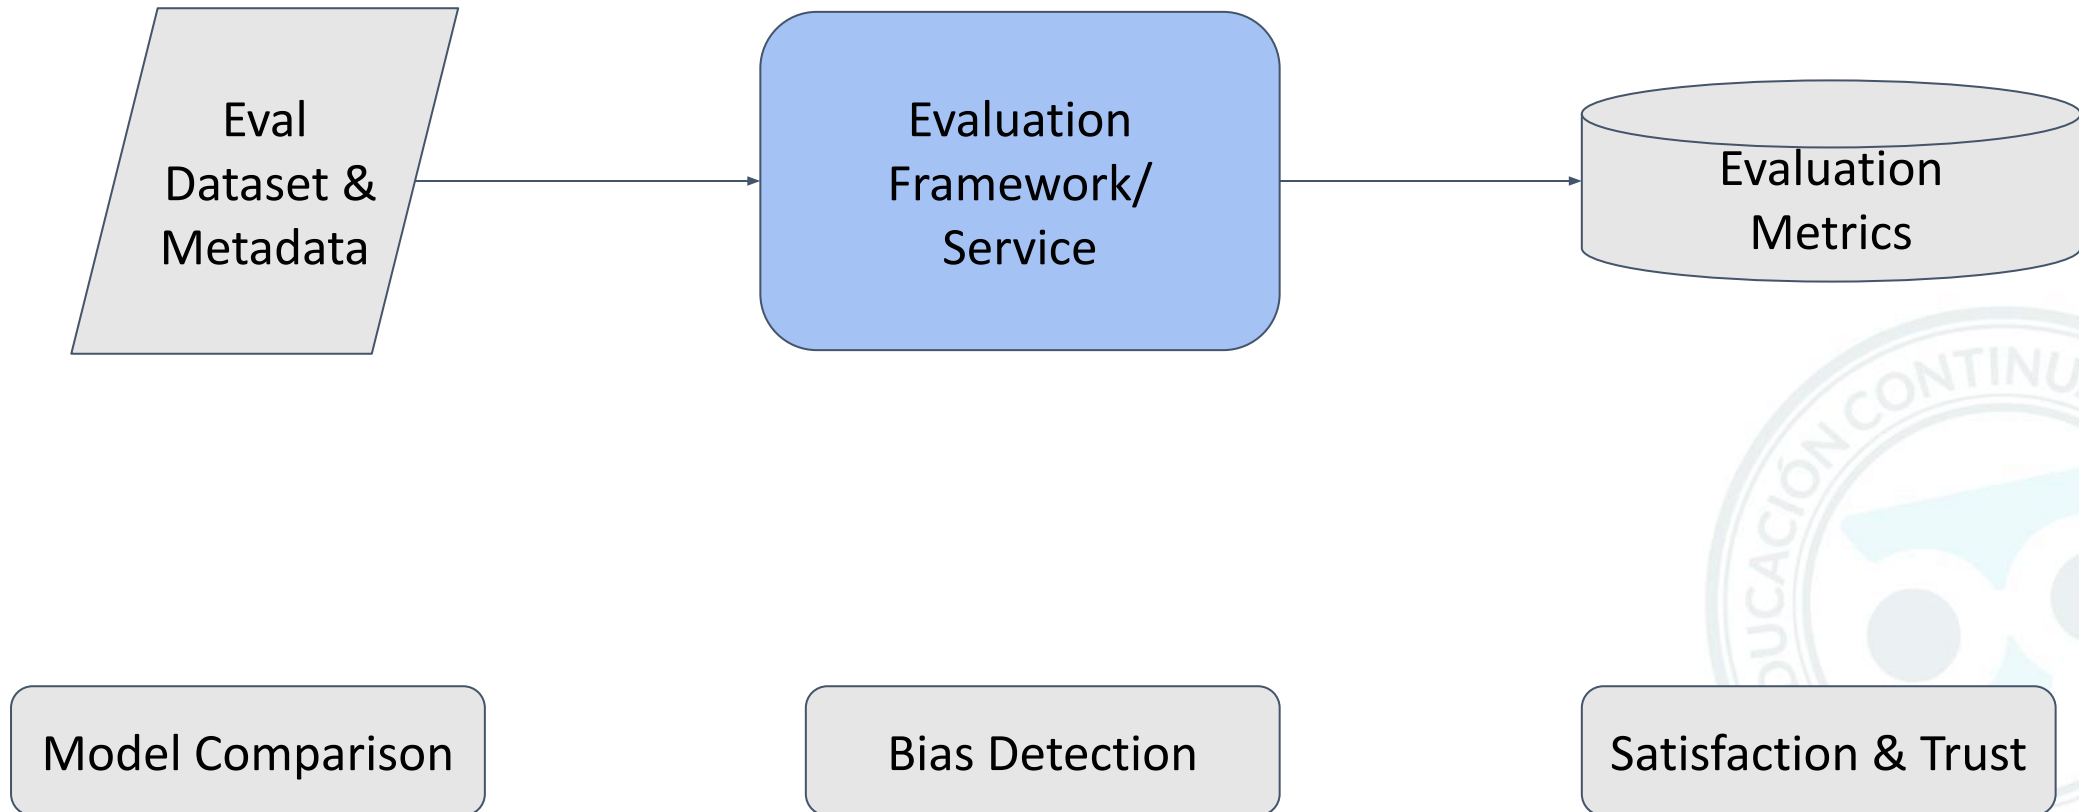
LLM evaluation of LLMs



- Prometheus
- JudgeLM



Evaluation - One of the Most Impactful Parts



Evaluating Gen AI systems

Key advice

Augment existing eval tooling

Why?

- Much tooling is reusable: MLflow, data pipelines, etc.
- New metrics can be added to existing systems

How?

- Adopt metrics from classic areas: toxicity (NLP), precision/recall (IR), ...
- Use new tools like LLM-as-a-judge
- Evaluate both the components + system as a whole

Build user feedback into your app

Why?

- Users can be the best judges
- Build proprietary datasets for the future fine-tuning and pretraining

How?

- Consider implicit and explicit feedback
- Manage feedback like any other data: same governance, same ETL, etc.



Evaluating Gen AI systems with mlflow

Batch evaluation in code

- LLM-as-a-judge
- Human evaluation using ground truth data
- New metrics for Gen AI, NLP, and retrieval

```
from mlflow.metrics.genai.metric_definitions import answer_relevance

answer_relevance_metric = answer_relevance(model="endpoints:/gpt-4")

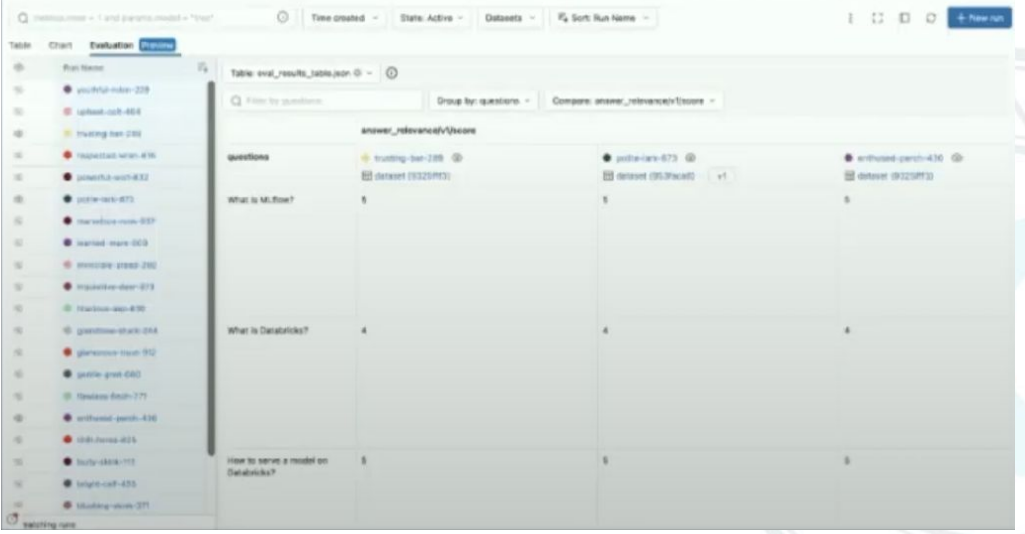
results = mlflow.evaluate(
    model,
    eval_df,
    model_type="question-answering",
    evaluators="default",
    predictions="result",
    extra_metrics=[answer_relevance_metric, mlflow.metrics.latency()],
    evaluator_config={
        "col_mapping": {
            "inputs": "questions",
            "context": "source_documents",
        }
    }
)

print(results.metrics)

results.tables["eval_results_table"]
```

Interactive evaluation in UI

- Compare Multiple models and prompts visually
- Iteratively test new queries during development



The screenshot shows the MLflow Evaluation UI. On the left, a list of runs is visible. The main panel displays a table titled 'eval_results_table.json' with columns for 'questions', 'gpt-4o-mini', 'gpt-4o', and 'gpt-4o-mini'. The table contains three rows of questions and their corresponding relevance scores for each model.

questions	gpt-4o-mini	gpt-4o	gpt-4o-mini
What is MLflow?	5	5	5
What is DataRobot?	4	4	4
How to serve a model on DataRobot?	5	5	5

Class Repo

<https://github.com/LGuillermoAngaritaG/llmops-class-eia>

Books:

<https://www.databricks.com/sites/default/files/2024-06/2023-10-EB-Big-Book-of-MLOps-2nd-Edition.pdf>

List of applications for LLMOps:

<https://github.com/tensorchord/Awesome-LLMOps>

Courses:

<https://github.com/mlabonne/llm-course>

¡TU OPINIÓN CUENTA!

Estudiante de Educación Continua, te invitamos a completar nuestra breve encuesta de satisfacción.

¡Ayúdanos a mejorar tu experiencia educativa!

